

Formale Methoden im Software Entwurf

Verifikation mit SPIN/ Verification with SPIN

Reiner Hähnle

29 Oktober 2018

Usage Scenario of PROMELA

1. **Model** the **essential** features of a system in PROMELA
 - ▶ **Abstract** away from complex (numerical) computations
 - ▶ Make use of **non-deterministic** choice of outcome
 - ▶ Replace unbounded datastructures with **finite** approximations
 - ▶ Assume **fair** process scheduler
2. **Select properties** that the PROMELA model must satisfy
 - ▶ Mutual exclusion for access to critical resources
 - ▶ Absence of deadlock
 - ▶ Absence of starvation
 - ▶ Event sequences (e.g., system responsiveness)
3. **Verify** that all possible runs of PROMELA model **satisfy** properties
 - ▶ Typically, need several **iterations** to get model and properties right
 - ▶ Failed verification attempts provide feedback via **counter examples**

SPIN: Previous Lecture vs. This Lecture

Previous lecture

SPIN presented as a PROMELA **simulator** (one run at a time)

This lecture

Intro to SPIN as a **model checker** (**all** runs & properties)

What Does A Model Checker Do?

Model Checker (MC) is designed to prove the designer wrong

MC does **not** try to prove correctness properties:
It tries the opposite!

MC tuned to **find counter example** to correctness property

Why can an MC also **prove** correctness properties?

Absence of any counter example proves stated correctness properties

MC's **search** for counter examples is **exhaustive**



“How often have I said to you that when you have eliminated the impossible, whatever remains, however improbable, must be the truth?”

What Does A Model Checker Do?

Model Checker (MC) is designed to prove the designer wrong

MC does **not** try to prove correctness properties:
It tries the opposite!

MC tuned to **find counter example** to correctness property

Why can an MC also **prove** correctness properties?

Absence of any counter example proves stated correctness properties

MC's **search** for counter examples is **exhaustive**

What Does “Exhaustive Search” Mean?

Exhaustive search

=

To resolve any non-determinism in all possible ways

For model checking PROMELA code,

two kinds of non-determinism need to be resolved:

► **Explicit, local:**

Overlapping guards in if/do statements

:: guardX -> ...

:: guardY -> ...

► **Implicit, global:**

Scheduling of concurrent processes
(see next lecture)

Model Checker for This Course: SPIN

SPIN: “Simple Promela Interpreter”

The name is a serious understatement!

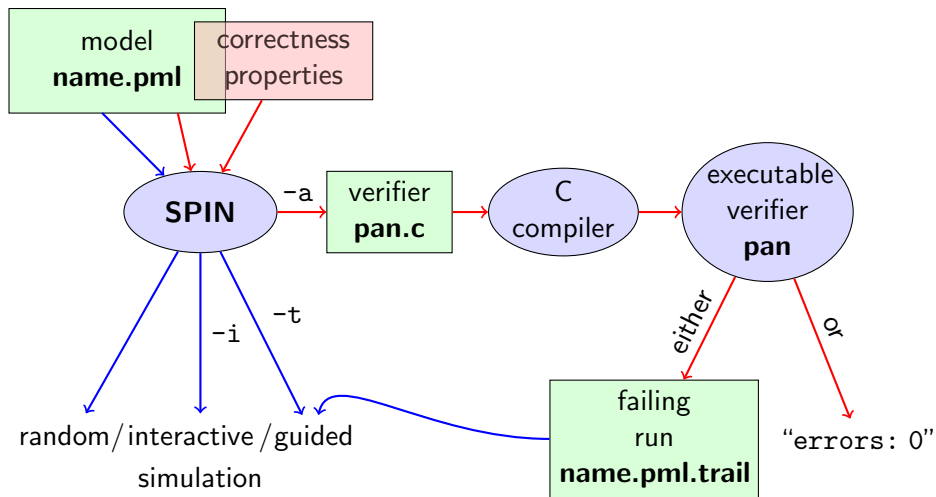
Main functionalities of SPIN:

1. Simulating a model (randomly/interactively/**guided**)
as seen in previous lecture
2. Generating a **verifier**

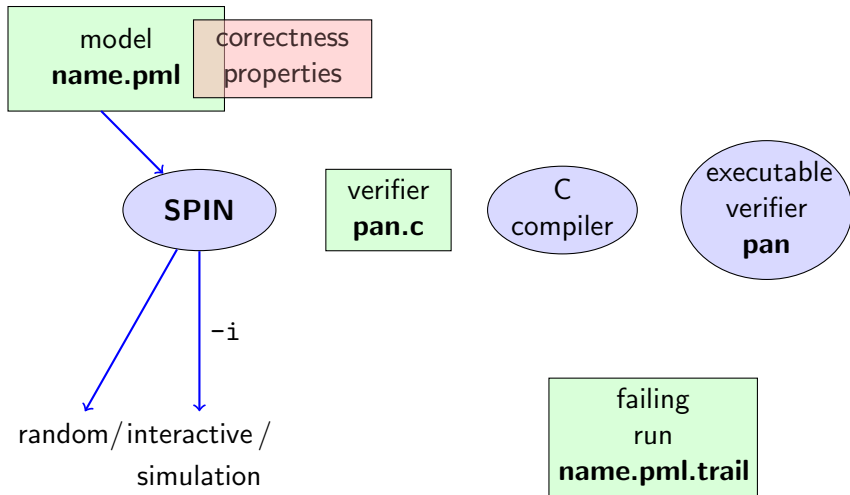
Verifier generated by SPIN is a C program performing **model checking**:

- ▶ Exhaustively **checks** PROMELA **model** against correctness properties
- ▶ In case the check is negative:
generates a **failing run** of the model, **to be simulated by SPIN**


SPIN Workflow: Overview, Refined



Plain Simulation with SPIN (Previous Lecture)



Model Checking with SPIN



Meaning of Correctness relative to Properties

Given PROMELA model M , and correctness properties C_1, \dots, C_n

Assume there is a check whether a PROMELA run R **satisfies** property C

Definition (Correctness of PROMELA model relative to property)


Let R_M be set of **all possible runs** of PROMELA model M

- ▶ For each correctness property C_i ,
 R_{M,C_i} is the set of all **runs** of M **satisfying** C_i (clearly, $R_{M,C_i} \subseteq R_M$)
- ▶ M is **correct relative to** C_1, \dots, C_n iff $R_M = (R_{M,C_1} \cap \dots \cap R_{M,C_n})$
- ▶ If M is not correct, then
each $r \in (R_M \setminus (R_{M,C_1} \cap \dots \cap R_{M,C_n}))$ is a **counter example**

We know how to write models M in PROMELA

But how to write **correctness properties**?

Stating Correctness Properties



Correctness properties can be stated **within** or **outside** a PROMELA model

Stating properties within PROMELA model

- ▶ **Assertion statements** (today)
- ▶ Meta labels ("Markierungen")
 - ▶ **end labels** (today)
 - ▶ accept labels
 - ▶ progress labels

Stating properties outside PROMELA model

- ▶ Never claims
- ▶ Temporal logic formulas

Assertion (“Zusicherung”) Statements

Definition (Assertion statement)

Assertion statements in PROMELA have the form

`assert(expr)`

where *expr* is any PROMELA expression.

Typically, (but not necessarily) *expr* is of type `bool`

`assert(expr)` in PROMELA can take any statement position:

```
...                               ...
stmt1;                           if
assert(max == a);                :: b1 -> stmt3;
stmt2;                           assert(x < y)
...                               :: b2 -> stmt4
                                ...
```

Meaning of **General** Assertion Statements

`assert(expr)`

- ▶ Has no effect if *expr* evaluates to **non-zero value**
- ▶ Triggers an error message if *expr* evaluates to **0**

This holds in both simulation and model checking mode

Recall:

`bool true false` is syntactic sugar for

`bit 1 0`

⇒ General case covers Boolean case

Instead of using “printf”s for Debugging ...

```
byte a, b, max;
select(a: 1 .. 3);
select(b: 1 .. 3);
if
  :: a >= b -> max = a
  :: a <= b -> max = b
fi;
printf("maximum of %d and %d is %d\n", a, b, max)
```

Command Line Execution

(simulate, check whether printed outcome is correct — oracle problem)

```
> spin [-i] max.pml
```


... we can employ **Assertions**

Quoting from file maxA.pml:


```
byte a, b, max;  
select(a: 1 .. 3);  
select(b: 1 .. 3);  
if  
  :: a >= b -> max = a  
  :: a <= b -> max = b  
fi;  
assert( max == (a>b -> a : b) )
```

First example with a formal **correctness property**

We can do **model checking**, for the first time!

(Historic moment in the course)

Model Checking, First Step: Generate Verifier in C



Command Line Execution


Generate Verifier in C

```
> spin -a maxA.pml
```

SPIN Generates **Verifier** in C, called **pan.c**

(plus auxiliary files)

Second Step: Compile To Executable Verifier



Command Line Execution


Compile to executable verifier

```
> gcc -o pan pan.c
```

C compiler generates **executable verifier pan**

pan: historically “**p**rotocol **a**nalyzer”, now “**p**rocess **a**nalyzer”

Third Step: Run Verifier (= Model Checking)




Command Line Execution

Run verifier pan

> ./pan or > pan (if current directory in search path)

- ▶ prints "errors: 0" (buried in a lot of extra info) ⇒ **Correctness**
Property verified!

Third Step: Run Verifier (= Model Checking)



Command Line Execution


Run verifier **pan**

> `./pan` or > `pan` (if current directory in search path)

- ▶ prints "errors: n " ($n > 0$) \Rightarrow counter example found!
records failing run in **maxF.pml.trail**

Errors Found: Guided Simulation

To **examine failing run**: employ **simulation mode**, “guided” by trail file.



Command Line Execution

Run verification of a faulty program, and then:

```
> spin -t -p -l maxF.pml
```

Output of Guided Simulation

Can look like:

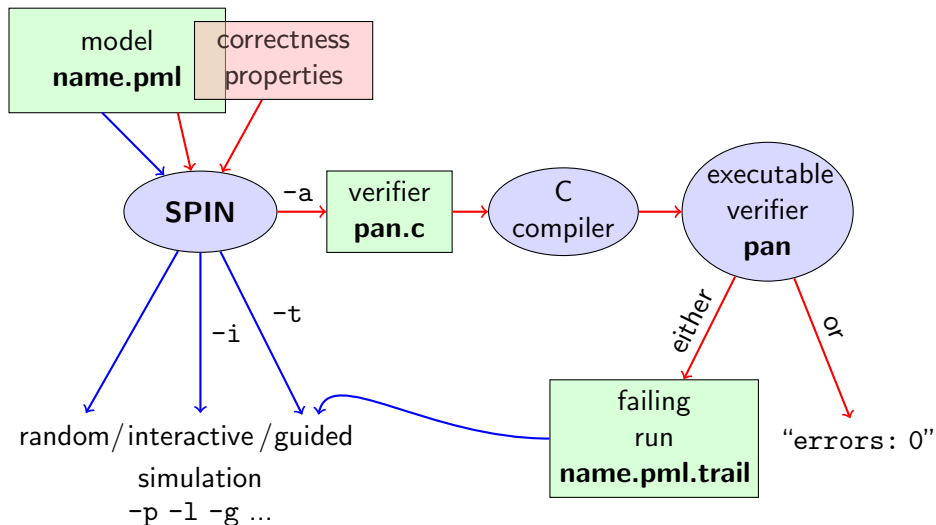
```
1: proc    0 (P) maxF.pml:7 (state 1)  [a = 1]
           P(0):a = 1
2: proc    0 (P) maxF.pml:12 (state 6) [b = 1]
           P(0):b = 1
3: proc    0 (P) maxF.pml:20 (state 13) [((a<=b))]
3: proc    0 (P) maxF.pml:20 (state 14) [max = b+1]
           P(0):max = 2
spin: maxF.pml:25, Error: assertion violated
spin: text of failed assertion:
      assert((max==( ((a>b)) -> (a) : (b) )))
```

Assignment statements executed in the run

Values of variables whenever updated

What did we do so far?

Following whole cycle (simple example, assertions only)



More Examples: Integer Division

```
int dividend = 15;
int divisor  = 4;
int quotient, remainder;

quotient = 0;
remainder = dividend;
do
    :: remainder > divisor -> quotient++;
    remainder = remainder - divisor
    :: else -> break
od;
printf("%d divided by %d = %d, remainder = %d\n",
       dividend, divisor, quotient, remainder)
```

- ▶ Simulate divide1, come up with assert in divide2
- ▶ With iSPIN: verify divide2, inspect trail, analyse, fix, verify

More Examples: Greatest Common Divisor

```
int x = 15, y = 20;
int a, b;
a = x; b = y;
do
  :: a > b -> a = a - b
  :: b > a -> b = b - a
  :: a == b -> break
od;
printf("The GCD of %d and %d = %d\n", x, y, a)
```

Full functional verification not possible here

- ▶ To express “greatest” property need quantification in expressions

Still, assertions can perform **sanity check**/partial verification (gcd.pml)

- ▶ Typical situation when employing model checking

Typical Command Lines

Some patterns for frequently used command line sequences:

Random simulation

```
spin name.pml or spin -v -l name.pml
```

Interactive simulation

```
spin -i name.pml
```

Model checking

```
spin -a name.pml  
gcc -o pan pan.c  
./pan | grep errors
```

and in case of errors:

```
spin -t -p -l -g name.pml
```

... or use ISPIN!

generates commands automatically

Ben-Ari produced **Spin Reference Card**, summarizing

- ▶ Typical command line sequences
- ▶ Options for
 - ▶ SPIN
 - ▶ gcc
 - ▶ pan
- ▶ PROMELA
 - ▶ datatypes
 - ▶ operators
 - ▶ statements
 - ▶ guarded commands
 - ▶ processes
 - ▶ channels
- ▶ Temporal logic syntax

Why SPIN?

Industrial-strength, “mainstream” software model checking

- ▶ SPIN targets **software**, instead of hardware verification
- ▶ 2001 ACM Software Systems Award (other winning software systems include: Unix, TCP/IP, WWW, Tcl/Tk, Java)
- ▶ Used for safety critical applications
- ▶ Annual SPIN user workshops series held since 1995
- ▶ Based on standard theory of ω -automata and linear temporal logic

Why SPIN? (Cont'd)

Suitable for Teaching

- ▶ PROMELA and SPIN are rather simple to use
- ▶ Distributed **freely** as research tool, well-documented actively maintained, large user-base in academia and in industry
- ▶ Good **portability**, quite easy to install, local installation ok
- ▶ Understand one representative system well, rather than many systems superficially
- ▶ Availability of good course **book** (Ben-Ari)
- ▶ Availability of **GUI**-based front end

Catching A Different Type of Error

Look again at max2.pml:

```
byte a, b, max;
select(a: 1 .. 3);
select(b: 1 .. 3);
if
  :: a >= b -> max = a;
  :: b <= a -> max = b;
fi;
printf("maximum of %d and %d is %d\n", a, b, max)
```

Simulate a few times

⇒ strange occasional “timeout” message

Generate and execute verifier **pan**

⇒ reports “errors: 1”

What is going on here?

Catching A Different Type of Error

Further inspection of **pan** output:

```
...  
pan: invalid end state (at depth 1)  
pan: wrote max2.pml.trail  
...
```


Legal and Illegal Blocking

There is no guard that covers $a < b \Rightarrow$ process blocks!

A process may legally block, as long as some other process can proceed

Blocking for letting others proceed is useful, and typical,
for concurrent and distributed models (for example, protocols)

But it is an error if a process blocks while no other process can proceed

\Rightarrow “Deadlock”

In `max2.pml` there exist runs when no process can proceed

Valid End States

Definition (Valid End State)

An **end state** of a **run** is **valid** iff the location counter of **each process** is at an **end location**.

Definition (End Location)

End locations of a process P are:

- ▶ P's textual end
- ▶ Any location marked with an **end label**: "endxxx:"

End labels not useful in **max2.pml**, but elsewhere, they are

Example: `verify end.pml` (without/with end labels)

Checking for invalid end states can be disabled in `pan` (option `-E`)

There are runs without any end state (e.g., non-terminating loop)

Literature for this Lecture

Ben-Ari Chapter 2
Section 4.7.1
Section 4.7.2